

MovieLens Recommendation system Project

Matias Garcia Mazzaro

octubre 22, 2020

Contents

The Project	2
The Objective	3
1 Download and construct the dataset	4
2 Analysis	6
2.1 Data partition	6
2.1.1 Exploring the data before to start	6
2.2 Creation of the RMSE function	7
2.3 The simplest model	8
2.4 <i>Modeling movie effect</i>	9
2.5 <i>Modeling user effect</i>	10
2.6 Regularization models	11
3 Improve it.	14
3.1 Data transformation	14
3.1.1 Dates transformation	14
3.1.2 Genres transformation	16
3.2 Adding the Oldness Movie effect	18
3.3 Moving the average in a function	19
3.4 Add the genre to the model	21
4 Validation	26
Results	28
Conlusions	29

The Project

This is a Machine Learning Project named MovieLens, it is a recommendation system of movies based in 10 millions dataset. The source of the dataset in this work is provided for [Group Lens](#), for download, press [here!!](#)

This work is part of the final project for the Data Science Professional certificate of HarvardX. In October 2006, Netflix offered a challenge to the data science community: improve our recommendation algorithm by 10% and win a million dollars.

The Objective

We will explore the data and construct a Machine Learning algorithm, describing the complete process, the method to compare different models will be through RMSE (root mean squared error) between the true data and predicted data.

We need create a data set that we will work and a validation data set. But we need to be sure don't use the validation dataset for define the model. For this question, we will create a data partition of the training data for test it and define the model.

We need obtain a $RMSE < 0.86490$ in the final validation for the total calibration.

Chapter 1

Download and construct the dataset

The Netflix data is not publicly available, but the GroupLens research lab generated their own database with over 20 million ratings for over 27,000 movies by more than 138,000 users.

We start with the creation of **edx** and **validation** objects and loading the packages and libraries we will use.

```
if(!require(tidyverse)) install.packages("tidyverse", repos = "http://cran.us.r-project.org")
if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")
if(!require(data.table)) install.packages("data.table", repos = "http://cran.us.r-project.org")

library(tidyverse)
library(caret)
library(data.table)

dl <- tempfile()
download.file("http://files.grouplens.org/datasets/movielens/ml-10m.zip", dl)

ratings <- fread(text = gsub("::", "\t", readLines(unzip(dl, "ml-10M100K/ratings.dat"))),
                 col.names = c("userId", "movieId", "rating", "timestamp"))

movies <- str_split_fixed(readLines(unzip(dl, "ml-10M100K/movies.dat")), "\\::", 3)
colnames(movies) <- c("movieId", "title", "genres")

# if using R 3.6 or earlier:
movies <- as.data.frame(movies) %>% mutate(movieId = as.numeric(levels(movieId))[movieId],
                                           title = as.character(title),
                                           genres = as.character(genres))

# if using R 4.0 or later:
movies <- as.data.frame(movies) %>% mutate(movieId = as.numeric(movieId),
                                           title = as.character(title),
                                           genres = as.character(genres))
```

```

movielens <- left_join(ratings, movies, by = "movieId")

# Validation set will be 10% of MovieLens data
set.seed(1, sample.kind="Rounding") # if using R 3.5 or earlier, use 'set.seed(1)'
test_index <- createDataPartition(y = movielens$rating, times = 1, p = 0.1, list = FALSE)
edx <- movielens[-test_index,]
temp <- movielens[test_index,]

# Make sure userId and movieId in validation set are also in edx set
validation <- temp %>%
  semi_join(edx, by = "movieId") %>%
  semi_join(edx, by = "userId")

# Add rows removed from validation set back into edx set
removed <- anti_join(temp, validation)
edx <- rbind(edx, removed)

rm(dl, ratings, movies, test_index, temp, movielens, removed)

```

Chapter 2

Analysis

It's important don't use the the **validation** data for train and test, it is because we need a efficient model for the data in the future, the final objetive is recomend movies to users, the true rating will occur after our recommendation. The first step, is to take **edx** object and make a data partition **train_set** and **test_set**:

2.1 Data partition

```
test_index <- createDataPartition(y = edx$rating, times = 1, p = 0.2,  
                                  list = FALSE)  
train_set <- edx[-test_index,]  
test_set <- edx[test_index,]
```

for be sure there are the same users and movies in the test set and training set;

```
test_set <- test_set %>%  
  semi_join(train_set, by = "movieId") %>%  
  semi_join(train_set, by = "userId")
```

2.1.1 Exploring the data before to start

For explore the dataset, we start for know how many diferents movies are in the dataset, and how many diferents users.

```
train_set %>%  
  summarize(n_users = n_distinct(userId),  
            n_movies = n_distinct(movieId))
```

```
##   n_users n_movies
## 1   69878   10641
```

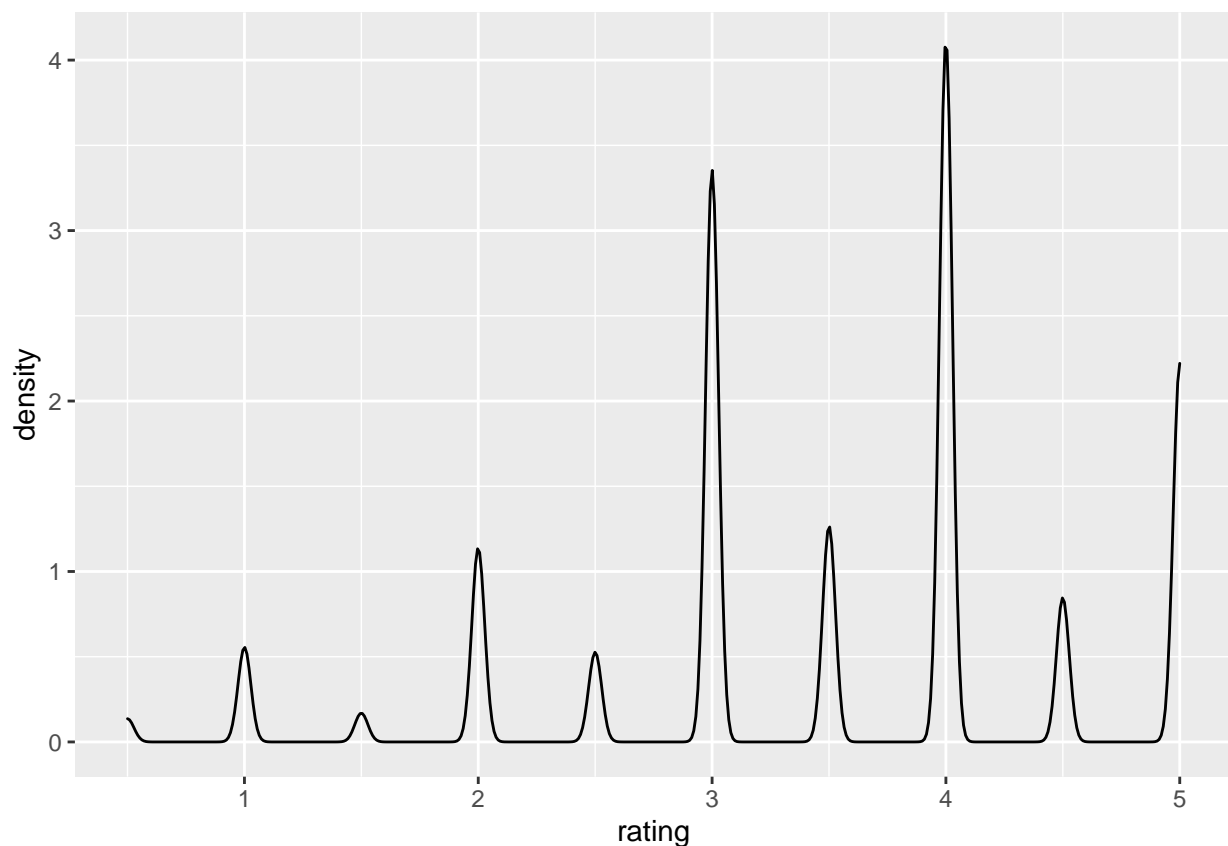
Also we can observe that the “genres” variable have differents asociated genres.

```
head(train_set$genres)
```

```
## [1] "Action|Crime|Thriller"      "Action|Drama|Sci-Fi|Thriller"
## [3] "Action|Adventure|Sci-Fi"    "Action|Adventure|Drama|Sci-Fi"
## [5] "Children|Comedy|Fantasy"    "Comedy|Drama|Romance|War"
```

We can observe also the distribution density of ratings.

```
edx %>% group_by(rating) %>% ggplot(aes(rating)) + geom_density()
```



2.2 Creation of the RMSE function

With the porpouse of messure the *root mean squared error* (RMSE), we define;

$$RMSE = \sqrt{\frac{1}{N} \sum_{u,i} (\hat{y}_{u,i} - y_{u,i})^2}$$

Where y is the rating for each movie i by user u and our prediction, and N being the number of user/movie combinations and the sum occurring over all these combinations.

```
RMSE <- function(true_ratings, predicted_ratings){
  sqrt(mean((true_ratings - predicted_ratings)^2))
}
```

The RMSE is similarly to a standard deviation: it is the typical error we make when predicting a movie rating. If this number is larger than 1, it means our typical error is larger than one star, which is not good.

2.3 The simplest model

The simplest possible recommendation system is the same rating for all movies regardless of user. A model that assumes the same rating for all movies and users with all the differences explained by random variation would look like this:

$$y_{u,i} = \mu + \epsilon_{u,i}$$

With ϵ independent errors sampled from the same distribution centered at 0 and μ the “true” rating for all movies.

```
mu_hat <- mean(train_set$rating)
mu_hat
```

```
## [1] 3.512574
```

```
naive_rmse <- RMSE(test_set$rating, mu_hat)
naive_rmse
```

```
## [1] 1.060704
```

We save the result in a new object for compare results later.

```
rmse_results <- tibble(method = "Just the average", RMSE = naive_rmse)
```

2.4 Modeling movie effect

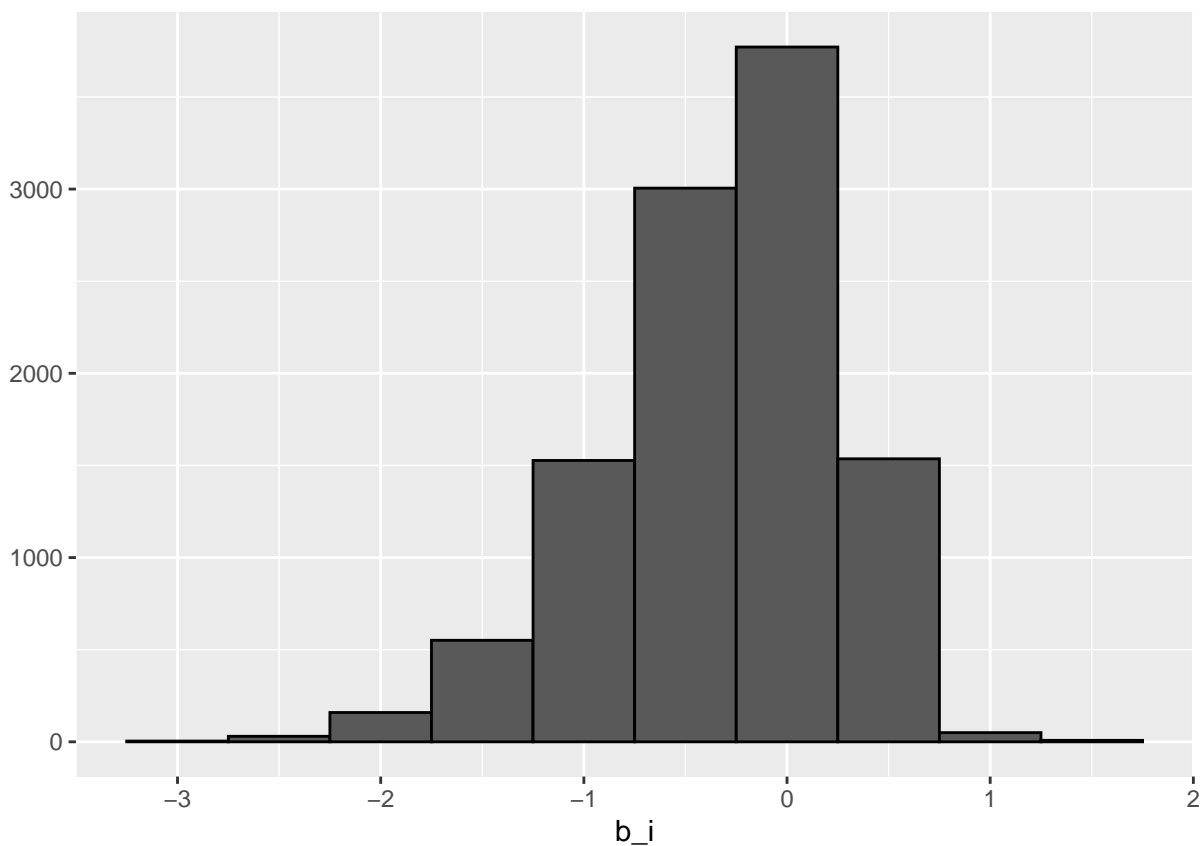
For model the *movie effect*, which implies, a *bias* added to the before model who move the mean for the mean of movie, we write;

$$y_{u,i} = \mu + b_i + \epsilon_{u,i}$$

```
mu <- mean(train_set$rating)

movie_avgs <- train_set %>%
  group_by(movieId) %>%
  summarize(b_i = mean(rating - mu))
```

The distribution of bias is;



We predict and test;

```
predicted_ratings <- mu + test_set %>%
  left_join(movie_avgs, by='movieId') %>%
  pull(b_i)
```

```
movie_effect <- RMSE(predicted_ratings, test_set$rating)
movie_effect
```

```
## [1] 0.9437144
```

We can see, the model including movie effect is better.

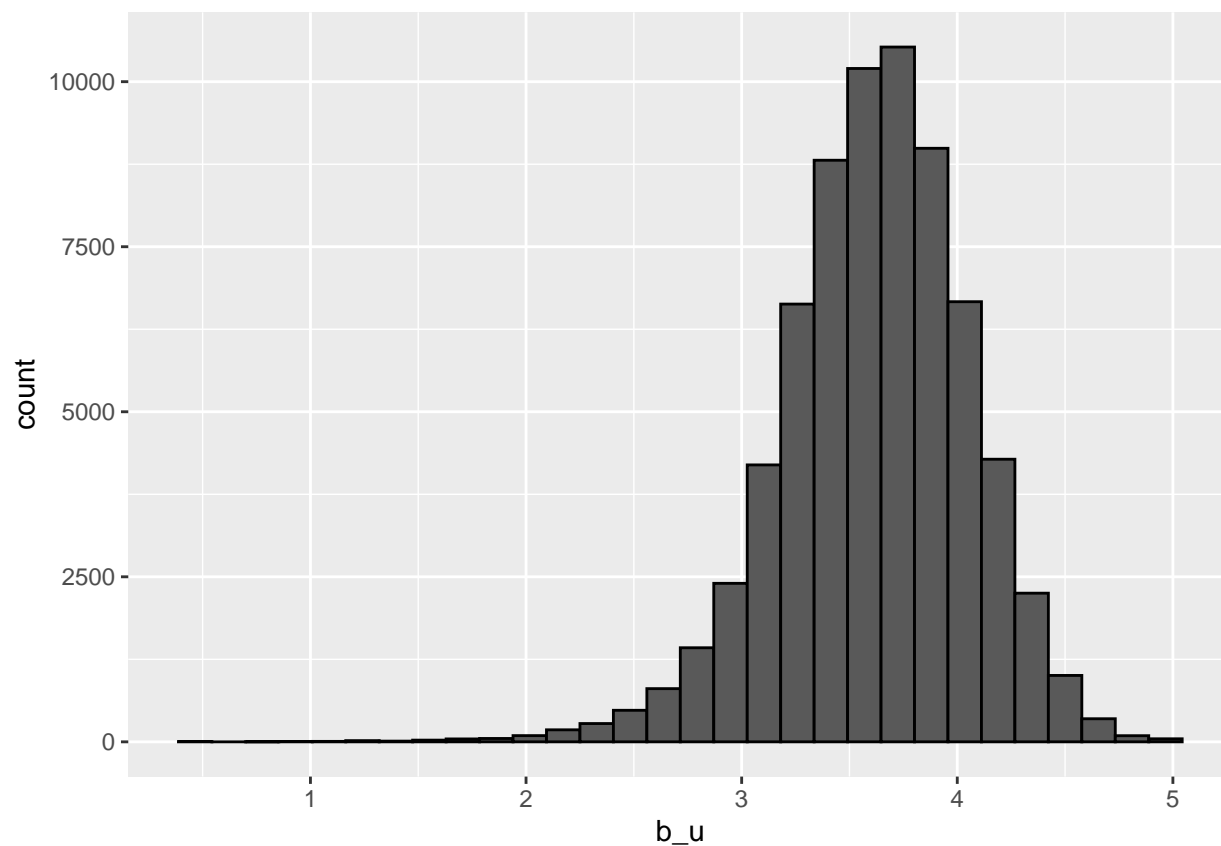
2.5 *Modeling user effect*

We add to the before model the user effect:

$$y_{u,i} = \mu + b_i + b_u + \epsilon_{u,i}$$

We observe the distribution for diferents users;

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```



We calculate de average for users;

```

user_avgs <- train_set %>%
  left_join(movie_avgs, by='movieId') %>%
  group_by(userId) %>%
  summarize(b_u = mean(rating - mu - b_i))

```

And predict movie effect adding user effect.

```

predicted_ratings <- test_set %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  mutate(pred = mu + b_i + b_u) %>%
  pull(pred)

user_effect <- RMSE(predicted_ratings, test_set$rating)
user_effect

```

```
## [1] 0.8661625
```

Sustancially diferent with others models.

2.6 Regularization models

If we explore the data for observe mistakes, we can see easily, there are movies with a fews ratings. Regularization permits us to penalize large estimates that are formed using small sample sizes.

For that we test which Lambda is the optimal ones that aport the minimal RMSE;

```

lambdas <- seq(0, 10, 0.25)

rmsees <- sapply(lambdas, function(l){

  mu <- mean(train_set$rating)

  b_i <- train_set %>%
    group_by(movieId) %>%
    summarize(b_i = sum(rating - mu)/(n()+1))

  b_u <- train_set %>%
    left_join(b_i, by="movieId") %>%
    group_by(userId) %>%
    summarize(b_u = sum(rating - b_i - mu)/(n()+1))

```

```

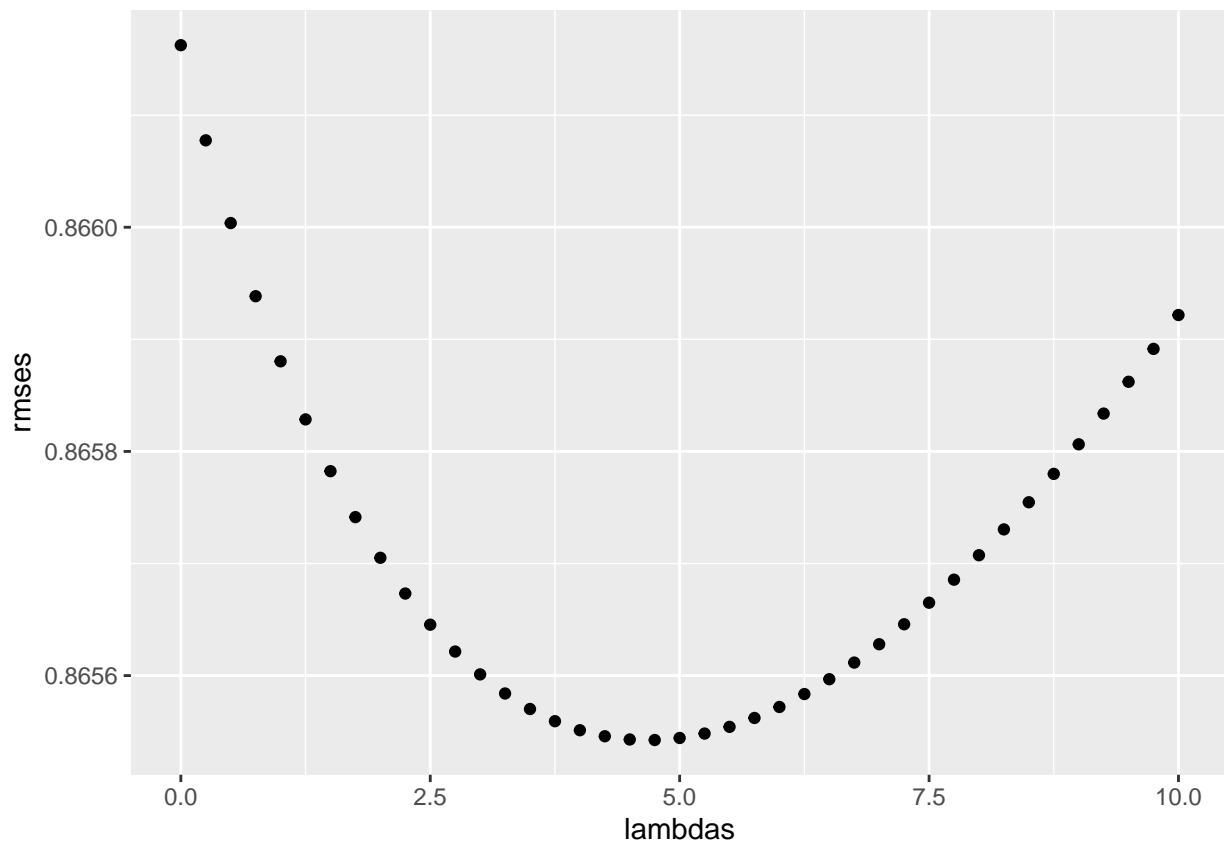
predicted_ratings <-
  test_set %>%
  left_join(b_i, by = "movieId") %>%
  left_join(b_u, by = "userId") %>%
  mutate(pred = mu + b_i + b_u) %>%
  pull(pred)

return(RMSE(predicted_ratings, test_set$rating))
})

```

We plot the Lambdas & the rmse for see how this affect, then we print the Lambda for the minimum RMSE.

```
qplot(lambdas, rmse)
```



```

lambda_opt <- lambdas[which.min(rmse)]
lambda_opt

```

```
## [1] 4.75
```

```
regularized_movie_user <- min(rmses)
```

The optimal Lambda is 4.75 .

```
min(rmses)
```

```
## [1] 0.8655425
```

And the rmse for Lambda 4.75 is 0.86524.

So, now we have this different RMSEs

```
## # A tibble: 4 x 2
##   method          RMSE
##   <chr>          <chr>
## 1 Just the average 1.06070
## 2 Movie effect Model 0.94371
## 3 Movie + User effect Model 0.86616
## 4 Regularized Movie + User effect Model 0.86554
```

Chapter 3

Improve it.

If we want to do this better, we need to think about the data. We observe, the date where the movie was filmed and the date where was ranked for each users. Another information we can consider is the gendres of the movies. For all this we need to process original data.

3.1 Data transformation

3.1.1 Dates transformation

If we look inside the data, we can observe all the titles of the films have the year were was released.

```
head(train_set$title)
```

```
## [1] "Net, The (1995)"          "Outbreak (1995)"
## [3] "Stargate (1994)"         "Star Trek: Generations (1994)"
## [5] "Flintstones, The (1994)" "Forrest Gump (1994)"
```

The first step is to add a column with the extraction of this year and count the oldness of the film, and then check it;

```
year_movie <-as.numeric(str_sub(train_set$title, start = -5, end = -2))
train_set <- train_set %>% mutate(year = year_movie)

year_movie <-as.numeric(str_sub(test_set$title, start = -5, end = -2))
test_set <- test_set %>% mutate(year = year_movie)

year_movie <-as.numeric(str_sub(validation$title, start = -5, end = -2))
validation <- validation %>% mutate(year = year_movie)
```

```

year_movie <- as.numeric(str_sub(edx$title, start = -5, end = -2))
edx <- edx %>% mutate(year = year_movie)

validation <- validation %>% mutate(Age = 2020 - year)
train_set <- train_set %>% mutate(Age = 2020 - year)
test_set <- test_set %>% mutate(Age = 2020 - year)
edx <- edx %>% mutate(Age = 2020 - year)

head(train_set)

```

```

##      userId movieId rating timestamp                title
## 1:      1      185      5 838983525             Net, The (1995)
## 2:      1      292      5 838983421             Outbreak (1995)
## 3:      1      316      5 838983392             Stargate (1994)
## 4:      1      329      5 838983392 Star Trek: Generations (1994)
## 5:      1      355      5 838984474   Flintstones, The (1994)
## 6:      1      356      5 838983653   Forrest Gump (1994)
##                                genres year Age
## 1:      Action|Crime|Thriller 1995  25
## 2: Action|Drama|Sci-Fi|Thriller 1995  25
## 3:      Action|Adventure|Sci-Fi 1994  26
## 4: Action|Adventure|Drama|Sci-Fi 1994  26
## 5:      Children|Comedy|Fantasy 1994  26
## 6:      Comedy|Drama|Romance|War 1994  26

```

Also we need to convert the timestamp in a year of rating.

```

class(validation$timestamp)

## [1] "integer"

validation <- validation %>% mutate(timestamp = as.Date.POSIXct(timestamp))
validation <- validation %>% mutate(year_rating = year(timestamp))

train_set <- train_set %>% mutate(timestamp = as.Date.POSIXct(timestamp))
train_set <- train_set %>% mutate(year_rating = year(timestamp))

test_set <- test_set %>% mutate(timestamp = as.Date.POSIXct(timestamp))
test_set <- test_set %>% mutate(year_rating = year(timestamp))

edx <- edx %>% mutate(timestamp = as.Date.POSIXct(timestamp))
edx <- edx %>% mutate(year_rating = year(timestamp))

head(validation)

```



```
##      userId movieId rating  timestamp
## 1:      1      231      5 1996-08-02
## 2:      1      480      5 1996-08-02
## 3:      1      586      5 1996-08-02
## 4:      2      151      3 1997-07-07
## 5:      2      858      2 1997-07-07
## 6:      2     1544      3 1997-07-07
##
##                                     title
## 1:                                Dumb & Dumber (1994)
## 2:                                Jurassic Park (1993)
## 3:                                Home Alone (1990)
## 4:                                Rob Roy (1995)
## 5:                                Godfather, The (1972)
## 6: Lost World: Jurassic Park, The (Jurassic Park 2) (1997)
##
##                                genres year Age year_rating
## 1:                                Comedy 1994  26          1996
## 2:      Action|Adventure|Sci-Fi|Thriller 1993  27          1996
## 3:                                Children|Comedy 1990  30          1996
## 4:      Action|Drama|Romance|War 1995  25          1997
## 5:                                Crime|Drama 1972  48          1997
## 6: Action|Adventure|Horror|Sci-Fi|Thriller 1997  23          1997
```

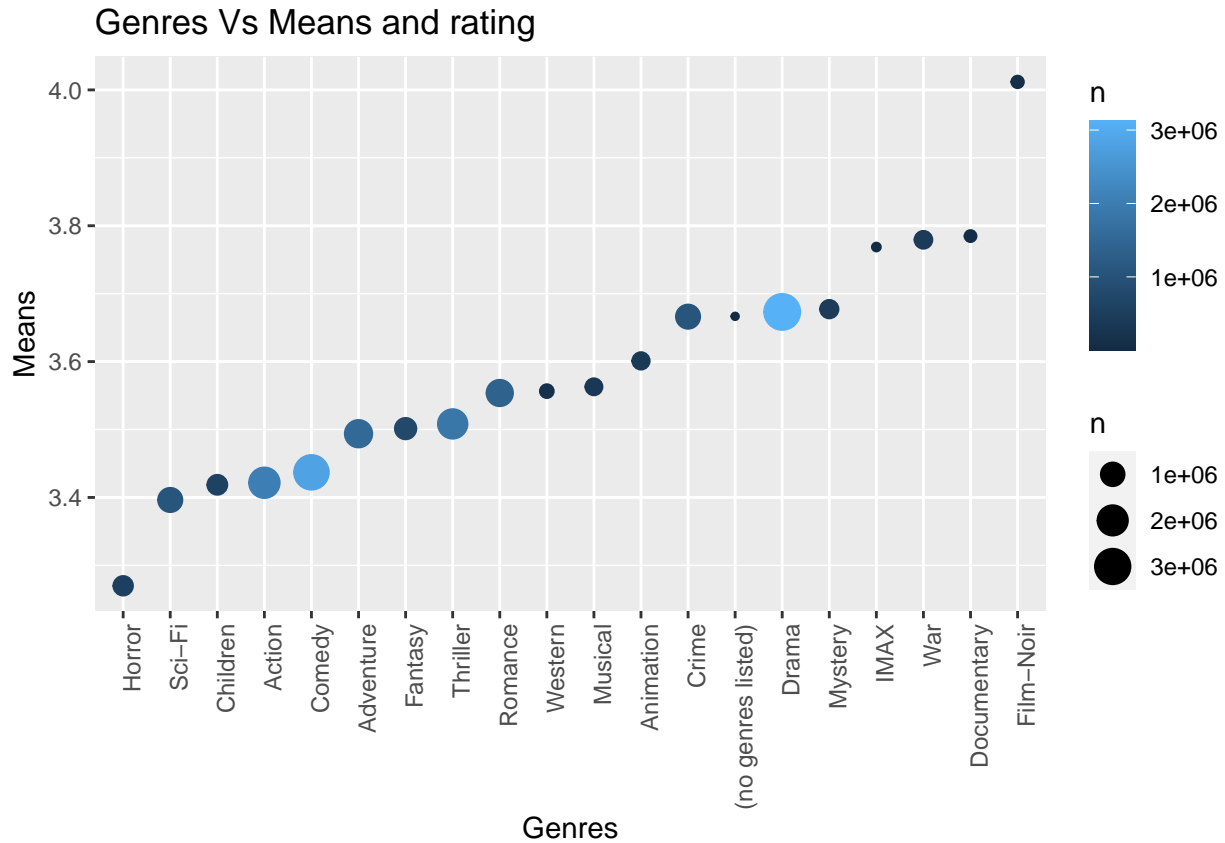
3.1.2 Genres transformation

We start we separate the genres and looks;

```
genres <- train_set %>% separate_rows(genres, sep = "\\|")
table_genres <- genres %>% group_by(genres) %>% summarize(n = n(), mean = mean(rating))

## 'summarise()' ungrouping output (override with '.groups' argument)

table_genres %>% mutate(genres = reorder(genres, mean)) %>% ggplot(aes(genres, mean, size
```



We can observe how diferents genres affect the mean ratings.

```
table_genres %>% mutate(dif = mean - mean(mean))
```

```
## # A tibble: 20 x 4
##   genres          n mean    dif
##   <chr>      <int> <dbl>  <dbl>
## 1 Drama    3127327  3.67  0.0857
## 2 Comedy   2832664  3.44 -0.150
## 3 Action   2048512  3.42 -0.166
## 4 Thriller  1861564  3.51 -0.0792
## 5 Adventure 1526571  3.49 -0.0936
## 6 Romance  1370133  3.55 -0.0337
## 7 Sci-Fi   1072759  3.40 -0.191
## 8 Crime    1062930  3.67  0.0788
## 9 Fantasy   740721  3.50 -0.0861
## 10 Children 590400  3.42 -0.169
## 11 Horror   553645  3.27 -0.318
## 12 Mystery  455346  3.68  0.0898
## 13 War      408941  3.78  0.192
## 14 Animation 373847  3.60  0.0136
```

## 15 Musical	346386	3.56	-0.0246
## 16 Western	151313	3.56	-0.0310
## 17 Film-Noir	95136	4.01	0.424
## 18 Documentary	74503	3.78	0.197
## 19 IMAX	6549	3.77	0.181
## 20 (no genres listed)	6	3.67	0.0792

3.2 Adding the Oldness Movie effect

```
mu <- mean(train_set$rating)
```

```
b_i <- train_set %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+4.75))
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
b_u <- train_set %>%
  left_join(b_i, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - mu)/(n()+4.75))
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
b_age <- train_set %>%
  left_join(b_i, by="movieId") %>%
  left_join(b_u, by="userId") %>%
  group_by(Age) %>%
  summarize(b_age = mean(rating - b_i - b_u - mu)/(n()+4.75))
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
predicted_ratings <-
  test_set %>%
  left_join(b_i, by = "movieId") %>%
  left_join(b_u, by = "userId") %>%
  left_join(b_age, by = "Age") %>%
  mutate(pred = mu + b_i + b_u + b_age) %>%
  pull(pred)
```

```
regu_user_movie_age <- RMSE(predicted_ratings, test_set$rating)
```

and compare the results with the before models;

```
rmse_results <- tibble(method = c("Just the average", "Movie effect Model", "Movie + User effect Model", "Regularized Movie + User effect Model", "Regularized Movie + User + Age effect Model"), rmse = c(1.06070, 0.94371, 0.86616, 0.86554, 0.86554))
```

```
## # A tibble: 5 x 2
##   method          RMSE
##   <chr>          <dbl>
## 1 Just the average 1.06070
## 2 Movie effect Model 0.94371
## 3 Movie + User effect Model 0.86616
## 4 Regularized Movie + User effect Model 0.86554
## 5 Regularized Movie + User + Age effect Model 0.86554
```

There is not significant results by oldness movie.

3.3 Moving the average in a function

Like we optimize the parameter “Lambda”, we will move the mean near the average μ , is possible there a value who report a better result for the RMSE. We plot the μ vs. rmse.

```
mu <- mean(train_set$rating)+seq(-0.2, 0.7, 0.05)

rmse_results <- sapply(mu, function(mu){

  b_i <- train_set %>%
    group_by(movieId) %>%
    summarize(b_i = sum(rating - mu)/(n()+4.75))

  b_u <- train_set %>%
    left_join(b_i, by="movieId") %>%
    group_by(userId) %>%
    summarize(b_u = sum(rating - b_i - mu)/(n()+4.75))

  b_age <- train_set %>%
    left_join(b_i, by="movieId") %>%
    left_join(b_u, by="userId") %>%
    group_by(Age) %>%
    summarize(b_age = mean(rating - b_i - b_u - mu)/(n()+4.75))

  predicted_ratings <-
    test_set %>%
    left_join(b_i, by="movieId") %>%
    left_join(b_u, by="userId") %>%
    left_join(b_age, by="Age") %>%
    summarize(predicted_ratings = rating - b_i - b_u - b_age)
```

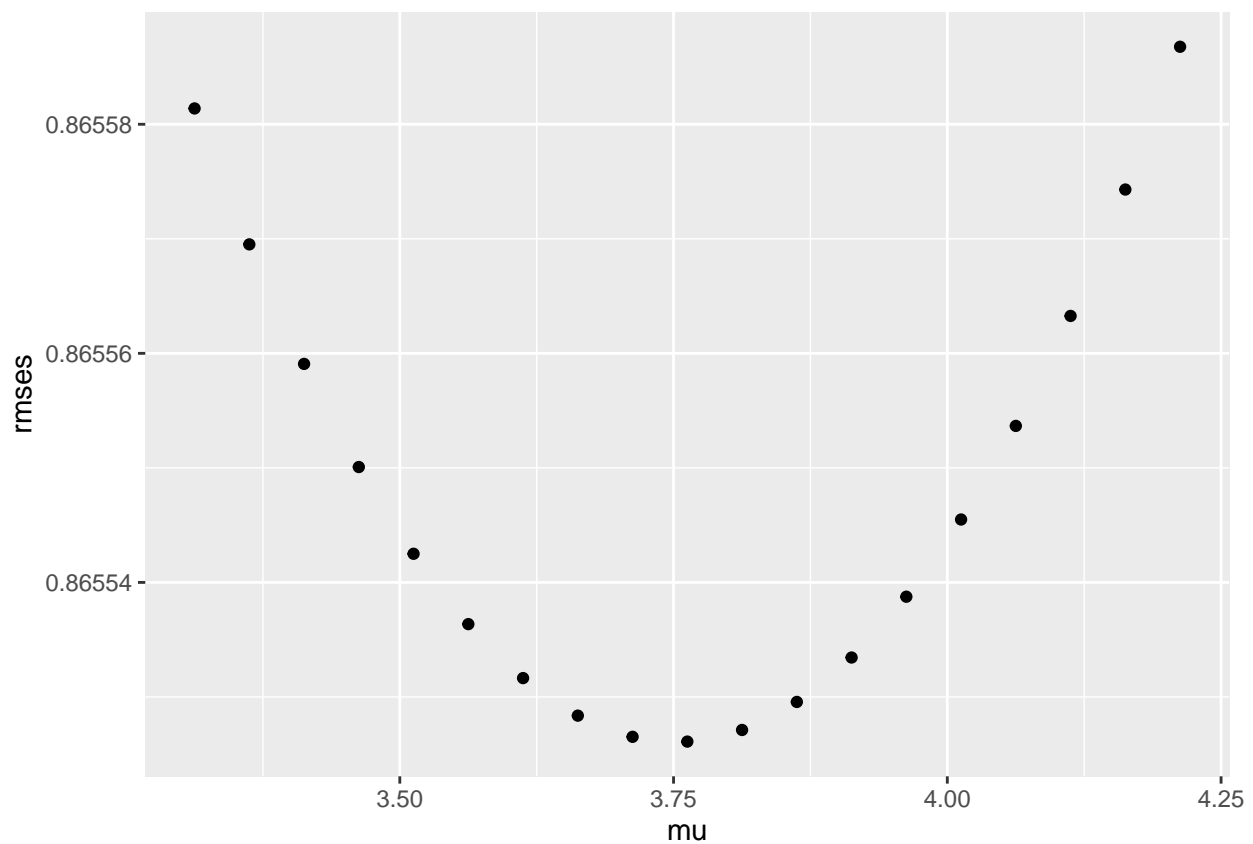
```

left_join(b_u, by = "userId") %>%
left_join(b_age, by = "Age") %>%
mutate(pred = mu + b_i + b_u + b_age) %>%
pull(pred)

return(RMSE(predicted_ratings, test_set$rating))
})

qplot(mu, rmses)

```



We can see there are a optimal μ over the average.

```

mu <- mu[which.min(rmses)]
mu

```

```
## [1] 3.762574
```

```
regu_user_movie_age_mu <- min(rmses)
```

We add this change in our table.

```
rmse_results <- tibble(method = c("Just the average", "Movie effect Model", "Movie + User effect Model"),
  rmse = c(1.06070, 0.94371, 0.86616, 0.86554, 0.86554, 0.86553))
```

```
## # A tibble: 6 x 2
##   method                RMSE
##   <chr>                 <dbl>
## 1 Just the average      1.06070
## 2 Movie effect Model    0.94371
## 3 Movie + User effect Model 0.86616
## 4 Regularized Movie + User effect Model 0.86554
## 5 Regularized Movie + User + Age effect Model 0.86554
## 6 Reg. Movie + User + Age effect Model (best mu) 0.86553
```

The new RMSE is better, but is not sufficient.

3.4 Add the genre to the model

We will try modeling also the genre, but, instead for each individual genres, for the combination of them. We can see there are near 800 different combinations.

```
train_set %>% group_by(genres) %>% summarize(b_gen = mu - mean(rating))
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
## # A tibble: 794 x 2
##   genres                b_gen
##   <chr>                 <dbl>
## 1 (no genres listed)    0.0959
## 2 Action                0.826
## 3 Action|Adventure      0.103
## 4 Action|Adventure|Animation|Children|Comedy -0.202
## 5 Action|Adventure|Animation|Children|Comedy|Fantasy 0.763
## 6 Action|Adventure|Animation|Children|Comedy|IMAX 0.526
## 7 Action|Adventure|Animation|Children|Comedy|Sci-Fi 0.694
## 8 Action|Adventure|Animation|Children|Fantasy 1.06
## 9 Action|Adventure|Animation|Children|Sci-Fi 0.785
## 10 Action|Adventure|Animation|Comedy|Drama 0.248
## # ... with 784 more rows
```

We training;

```

mu <- mean(train_set$rating)

b_i <- train_set %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+4.75))

## 'summarise()' ungrouping output (override with '.groups' argument)

b_u <- train_set %>%
  left_join(b_i, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - mu)/(n()+4.75))

## 'summarise()' ungrouping output (override with '.groups' argument)

b_age <- train_set %>%
  left_join(b_i, by="movieId") %>%
  left_join(b_u, by = "userId") %>%
  group_by(Age) %>%
  summarize(b_age = sum(rating - b_i - b_u - mu)/(n()+4.75))

## 'summarise()' ungrouping output (override with '.groups' argument)

b_gen <- train_set %>%
  left_join(b_i, by="movieId") %>%
  left_join(b_u, by = "userId") %>%
  left_join(b_age, by = "Age") %>%
  group_by(genres) %>%
  summarize(b_gen = sum(rating - b_i - b_u - b_age - mu)/(n()+4.75))

## 'summarise()' ungrouping output (override with '.groups' argument)

predicted_ratings <- test_set %>%
  left_join(b_i, by = "movieId") %>%
  left_join(b_u, by = "userId") %>%
  left_join(b_age, by = "Age") %>%
  left_join(b_gen, by = "genres") %>%
  mutate(pred = mu + b_i + b_u + b_age) %>%
  pull(pred)

regu_user_movie_age_gen <- RMSE(predicted_ratings, test_set$rating)

```

We explore the results.

```
rmse_results <- tibble(method = c("Just the average", "Movie effect Model", "Movie + User effect Model"),  
rmse_results
```

```
## # A tibble: 6 x 2  
##   method          RMSE  
##   <chr>          <chr>  
## 1 Just the average 1.06070  
## 2 Movie effect Model 0.94371  
## 3 Movie + User effect Model 0.86616  
## 4 Regularized Movie + User effect Model 0.86554  
## 5 Regularized Movie + User + Age effect Model 0.86554  
## 6 Reg. Movie + User + Age + Genre effect Model 0.86527
```

Now we obtain a significant result.

We fit again the μ parameter for more precision.

```
mu <- mean(train_set$rating)+seq(0.1, 0.3, 0.05)  
  
rmse_results <- sapply(mu, function(mu){  
  
  b_i <- train_set %>%  
    group_by(movieId) %>%  
    summarize(b_i = sum(rating - mu)/(n()+4.75))  
  
  b_u <- train_set %>%  
    left_join(b_i, by="movieId") %>%  
    group_by(userId) %>%  
    summarize(b_u = sum(rating - b_i - mu)/(n()+4.75))  
  
  b_age <- train_set %>%  
    left_join(b_i, by="movieId") %>%  
    left_join(b_u, by = "userId") %>%  
    group_by(Age) %>%  
    summarize(b_age = sum(rating - b_i - b_u - mu)/(n()+4.75))  
  
  b_gen <- train_set %>%  
    left_join(b_i, by="movieId") %>%  
    left_join(b_u, by = "userId") %>%  
    left_join(b_age, by = "Age") %>%  
    group_by(genres) %>%
```



```

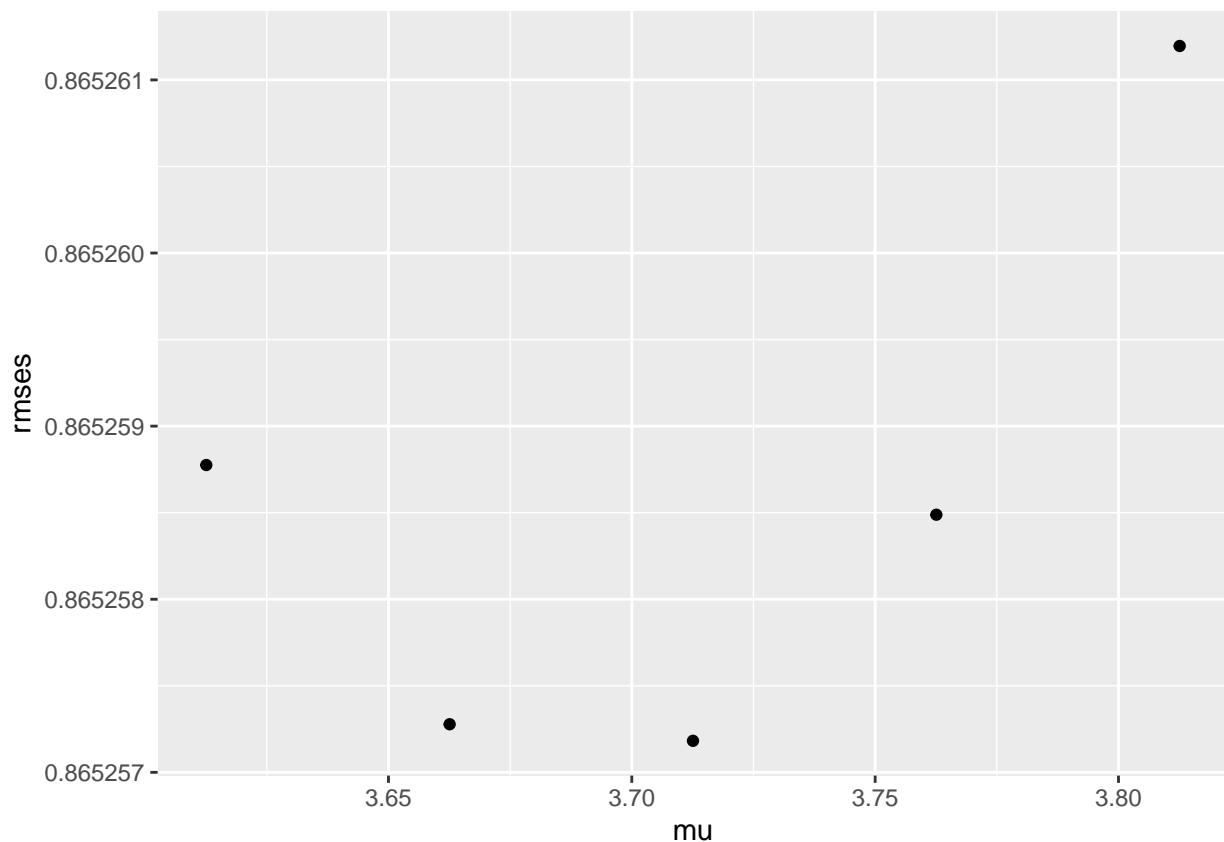
    summarize(b_gen = sum(rating - b_i - b_u - b_age - mu)/(n()+4.75))

predicted_ratings <- test_set %>%
  left_join(b_i, by = "movieId") %>%
  left_join(b_u, by = "userId") %>%
  left_join(b_age, by = "Age") %>%
  left_join(b_gen, by = "genres") %>%
  mutate(pred = mu + b_i + b_u + b_age) %>%
  pull(pred)

return(RMSE(predicted_ratings, test_set$rating))
})

qplot(mu, rmses)

```



```
mu[which.min(rmses)]
```

```
## [1] 3.712574
```

```

regu_user_movie_age_gen_mu <- min(rmses)

rmse_results <- tibble(method = c("Just the average", "Movie effect Model", "Movie + User effect Model", "Movie + User + Age effect Model", "Movie + User + Age + Genre effect Model", "Movie + User + Age + Genre effect Model (mu opt)", "Movie + User + Age + Genre effect Model (mu opt)"))
rmse_results

```

```

## # A tibble: 7 x 2
##   method                                RMSE
##   <chr>                                <chr>
## 1 Just the average                    1.06070
## 2 Movie effect Model                  0.94371
## 3 Movie + User effect Model           0.86616
## 4 Regularized Movie + User effect Model 0.86554
## 5 Regularized Movie + User + Age effect Model 0.86554
## 6 Reg. Movie + User + Age + Genre effect Model 0.86527
## 7 Reg. Movie + User + Age + Genre effect Model (mu opt) 0.86526

```

Chapter 4

Validation

For test our best model, the last one, we need to process now with the validation set. For train our model we will use the complete training set “edx”. Like edx have more data, can be posible a improvement respect our last training.

We use our optimized parameters.

```
mu <- 3.712482

b_i <- edx %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+4.75))

b_u <- edx %>%
  left_join(b_i, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - mu)/(n()+4.75))

b_age <- edx %>%
  left_join(b_i, by="movieId") %>%
  left_join(b_u, by = "userId") %>%
  group_by(Age) %>%
  summarize(b_age = sum(rating - b_i - b_u - mu)/(n()+4.75))

b_gen <- edx %>%
  left_join(b_i, by="movieId") %>%
  left_join(b_u, by = "userId") %>%
  left_join(b_age, by = "Age") %>%
  group_by(genres) %>%
  summarize(b_gen = sum(rating - b_i - b_u - b_age - mu)/(n()+4.75))

predicted_ratings <- validation %>%
```

```

left_join(b_i, by = "movieId") %>%
left_join(b_u, by = "userId") %>%
left_join(b_age, by = "Age") %>%
left_join(b_gen, by = "genres") %>%
mutate(pred = mu + b_i + b_u + b_age) %>%
pull(pred)

```

```

validation_rmse <- RMSE(predicted_ratings, validation$rating)
validation_rmse

```

```
## [1] 0.8645024
```

We obtain a RMSE of our prediction and true data, of 0.8645024.

```

rmse_results <- tibble(method = c("Just the average", "Movie effect Model", "Movie + User effect Model", "Regularized Movie + User effect Model", "Regularized Movie + User + Age effect Model", "Reg. Movie + User + Age + Genre effect Model", "Reg. Movie + User + Age + Genre effect Model (mu opt)", "final model validation"))
rmse_results

```

```

## # A tibble: 8 x 2
##   method                                RMSE
##   <chr>                                <chr>
## 1 Just the average                    1.06070
## 2 Movie effect Model                 0.94371
## 3 Movie + User effect Model          0.86616
## 4 Regularized Movie + User effect Model 0.86554
## 5 Regularized Movie + User + Age effect Model 0.86554
## 6 Reg. Movie + User + Age + Genre effect Model 0.86527
## 7 Reg. Movie + User + Age + Genre effect Model (mu opt) 0.86526
## 8 final model validation             0.86450

```

Results

In the final validation, were obtain a RMSE of **0.86450** , it is better of the objctive (0.86490).

Conlusions

In the Netflix challenge the winning score was $RMSE=0.8712$. Data exploration is a process, we start with a very simple model and advance in more complex models trough we were knowing the data.

I obtain the objetive RMSE with a computational power of a desktop computer, and always is posible move the average in the base of the model for optimize, denote a coarse fit. Its posible improve it, ensemble models, and moving parameters, also using models like Random Forest, Matriz Factorization, SVD and PCA.